

Machine Learning (CS 181):

20. Markov Decision Processes

David C. Parkes and Sasha Rush

Spring 2017

1 / 51

Contents

1 Introduction

2 Planning (finite horizon)

3 Planning (infinite horizon)

- Bellman equations
- Value Iteration
- Policy Iteration

4 Conclusion

2 / 51

Overview

Supervised learning

$$D = \{(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_n, \mathbf{y}_n)\}$$

Neural networks, Naive Bayes, SVMs, random forests, linear regression, ...


Unsupervised learning

$$D = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$$

K-means, HAC, Bayesian Networks, topic models, Gaussian mixture models, HMMs...


Learning to act: embodied agents

$$D = (s_1, a_1, r_1, s_2, a_2, r_2, \dots)$$



3 / 51

Markov Decision Process




An MDP is specified by (S, A, r, p) :

- $S = \{1, \dots, |S|\}$ states
- $A = \{1, \dots, |A|\}$ actions
- reward function $r(s, a) \in \mathbb{R}$, for all states s , all actions a
- transition model $p(s' | s, a)$, for all states s , actions a , next states s'

A policy π is a mapping from states to actions. Want to find 'rewarding' policies..

4 / 51


Application 1: Robots



- States: physical location, objects in environment
- Actions: move, pick-up, drop, ...
- Reward: +1 if pick up dirty clothes, -1 if break dish, ...
- Transition model: describe actuators and uncertain environment

5 / 51

Application 2: Game of Go



- States: board position
- Actions: move a piece
- Reward: +1 if win the game, 0 if draw, -1 if lose the game
- Transition model: rules of game, response of other player

6 / 51

AlphaGo vs. Lee Sedol




- AlphaGo (DeepMind) defeated Lee Sedol, 4-1 in March 2016, the top Go player in the world
- AlphaGo combines Monte-Carlo tree search with deep neural nets (trained by supervised learning), with reinforcement learning.
- Learns both a 'policy network' (which action to play in which state) and a 'value network' (estimate of value of an action under self-play).

'Mastering the game of Go with deep neural networks and tree search', Silver et al., Nature **529**:484–582 (2016)

7 / 51

Application 3: Customer Service Agent



- States: summary of conversation so far
- Actions: words to utter
- Reward: +1 if solve caller's problem, -1 if need to go to human, -10 if caller hangs up angry
- Transition model: effect of words on state, next words or action from caller.

8 / 51

Working with MDPs

An MDP is a general probabilistic framework, and can be utilized in many different scenarios.

- Planning:

- Full access to the MDP, compute an optimal policy.
- “How do I act in a known world?”

- Policy Evaluation:

- Full access to the MDP, compute the ‘value’ of a fixed policy.
- “How will this plan perform under uncertainty?”

- Reinforcement Learning (next lecture):

- Limited access to the MDP.
- “Can I learn to act in an uncertain world?”

9 / 51

Different Objective Criteria

- Sequence of $s_1, a_1, r_1, s_2, a_2, r_2, \dots$; discrete time t
- Finite horizon, $T \geq 1$ steps

$$\text{utility} = \sum_{t=1}^T r(s_t, a_t)$$


- Infinite horizon, discount factor $\gamma \in (0, 1]$

$$\text{utility} = r(s_1, a_1) + \gamma \cdot r(s_2, a_2) + \gamma^2 \cdot r(s_3, a_3) + \dots$$

(Long-run average, $\lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^{\infty} r(s_t, a_t)$ is another objective criterion.)

10 / 51

Running Illustration: MDP on Gridworld




S	Location of the grid (x_1, x_2)
A	Local movements $\leftarrow, \rightarrow, \uparrow, \downarrow$
$r : S \times A \mapsto \mathbb{R}$	Reward function, e.g. make it to goal
$p(s' s, a)$	Transition model, e.g. deterministic or slippages

11 / 51

Example Gridworld (perfect actuator)

Optimal policy:




- $r(s, a) = -0.04$ for all states, actions except $(4, 2), (4, 3)$
- Bounce off obstacles
- Stop when get to $(4, 2), (4, 3)$ ('episodic')
- Perfect actuator

12 / 51

Gridworld Example (perfect actuator)

Optimal policy:




- $r(s, a) = -0.04$ for all states, actions except $(4, 2), (4, 3)$
- Bounce off obstacles
- Stop when get to $(4, 2), (4, 3)$ ('episodic')
- ~~Perfect actuator~~ imperfect actuator (prob. 0.1 in direction 90° left, prob. 0.1 in direction 90° right)?

13 / 51

Example (imperfect actuator)

In this case, optimal policy becomes:



- $r(s, a) = -0.04$ for all states, actions except $(4, 2), (4, 3)$
- Bounce off obstacles
- Stop when get to $(4, 2), (4, 3)$ ('episodic')
- ~~Perfect actuator~~ imperfect actuator (prob. 0.1 in direction 90° left, prob. 0.1 in direction 90° right)?

14 / 51

1 Introduction

2 Planning (finite horizon)


3 Planning (infinite horizon)

- Bellman equations
- Value Iteration
- Policy Iteration

4 Conclusion

15 / 51

Warm-up: Expectimax



S. Yoon

- Build out a look-ahead tree to the decision horizon; max over actions, exp over next states.
- Solve from the leaves, backing-up the expectimax values.
- Problem: computation is exponential in horizon.
- May expand the same subtree multiple times. (e.g., s_1, a_1 and s_1, a_2 may lead to same state.)

16 / 51

Finite-Horizon Planning: Value iteration

A dynamic programming approach.

Let $V_{(t)}^*(s)$ denote the total value from state s under optimal policy with t -steps-to-go, $\pi_{(t)}^*(s)$ the optimal action with t -periods-to-go.

Base case (for all states s):

$$V_{(1)}^*(s) = \max_a r(s, a).$$

Inductive case (for all states s , time-to-go $t = 2, \dots, T$):

$$V_{(t)}^*(s) = \max_{a \in A} \left[r(s, a) + \sum_{s' \in S} p(s' | a, s) V_{(t-1)}^*(s') \right]$$

Work back from last period to present. Can read-off the optimal policy.

Let $L = \max \#$ states reachable from any state under any action.

Computational complexity is $O(|A| \cdot |S| \cdot L \cdot T)$.

17 / 51

Example: Value iteration

$$V_{(t)}^*(s) = \max_{a \in A} (r(s, a) + \sum_{s' \in S} p(s' | a, s) V_{(t-1)}^*(s'))$$

Simple 5-state, 2-action gridworld. Stop when get to states 1 or 5.

$r(s, a)$	10	-1	-1	-1	5	optimal policy?
-----------	----	----	----	----	---	-----------------

$V_{(1)}^*(s)$	10	-1	-1	-1	5
----------------	----	----	----	----	---

$V_{(2)}^*(s)$	10	9	-2	4	5
----------------	----	---	----	---	---

e.g., $9 = \max(-1 + 10, -1 - 1)$,
 $-2 = \max(-1 - 1, -1 - 1)$

$V_{(3)}^*(s)$	10	9	8	4	5
----------------	----	---	---	---	---

e.g., $8 = \max(-1 + 9, -1 + 4)$

$V_{(4)}^*(s)$	10	9	8	7	5
----------------	----	---	---	---	---

e.g., $7 = \max(-1 + 8, -1 + 5)$

18 / 51

1 Introduction

2 Planning (finite horizon)

3 Planning (infinite horizon)

- Bellman equations
- Value Iteration
- Policy Iteration

4 Conclusion

19 / 51

MDP Value function

Consider an infinite time horizon, and a stationary and deterministic policy $\pi(s) \in A$.

This is without loss of generality (for discounted objective criterion).

Definition (MDP value function)

The MDP value function of a policy π from state s is

$$V^\pi(s) = \mathbf{E} \left[\sum_{t=1}^{\infty} \gamma^{t-1} r(s_t, \pi(s_t)) \right]$$

where $s_1 \triangleq s$, and $s_{t+1} \sim p(s' | s_t, \pi(s_t))$.

20 / 51

Policy Evaluation

We can expand this MDP value function as:

$$V^\pi(s) = \underbrace{r(s, \pi(s))}_{\text{reward now}} + \gamma \underbrace{\sum_{s' \in S} p(s' | s, \pi(s)) V^\pi(s')}_{\text{expected, discounted future reward}} \quad (1)$$

Definition (Policy evaluation)

For a given policy π , infinite time horizon, and discounting, evaluate the MDP value function.

We can solve system of linear equations (1) in time $O(|S|^3)$ via Gaussian elimination.

21 / 51

Bellman equations

The planning problem for an MDP is:

$$\pi^* \in \arg \max_{\pi} V^\pi(s).$$

(exists a solution that is optimal for every state s).

Definition (Bellman equations)

For an optimal policy π^* , we have

$$V^*(s) = \max_{a \in A} \left[r(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V^*(s') \right], \quad \forall s \quad (2)$$

This system of (non-linear) equations capture the principle of optimality.
The value of an optimal policy = value of doing the right thing now, considering the value that comes from optimal 'continuation.'

22 / 51

Value iteration

The Bellman equations suggest the following approach to planning:

- Initialize: $V(s) = 0$, for all states s

- Update step ('Bellman operator'):

$$V'(s) \leftarrow \max_{a \in A} \left[r(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V(s') \right], \quad \forall s$$

update value function using one-step look-ahead

- $V \leftarrow V'$

Continue until converge, find the fixpoint. Can then read-off the optimal policy via (2).

Computation $O(|S| \cdot |A| \cdot L)$ per iteration, where $L = \max\#$ states reachable from any state under any action.

23 / 51

Convergence of Value Iteration

- Contraction property for update $x' \leftarrow f(x)$:

$$\|f(x) - f(y)\| < \|x - y\|, \quad \text{for all } x \neq y$$

- e.g., $x' \leftarrow f(x) = x/2$, fixpoint $x^* = f(x^*) \Leftrightarrow x^* = 0$

- contraction: $(2, 8), (1, 4), (1/2, 2), \dots$

- By contraction property:

- f has a unique fixpoint, else $\|f(x^*) - f(y^*)\| = \|x^* - y^*\|$
(violation of contraction)


- update converges to the fixpoint, consider $x \neq x^*$,

$$\|f(x) - x^*\| = \|f(x) - f(x^*)\| < \|x - x^*\|.$$

- The Bellman operator is a contraction when discount factor $\gamma < 1$,
and where $\|\mathbf{V}\| = \max_s |V(s)|$ ('max-norm')

24 / 51


Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

25 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

26 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

27 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

28 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

29 / 51


Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

30 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

31 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

32 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

33 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

34 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

35 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

36 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

37 / 51

Example: Value iteration in GridWorld



(D. Klein and P. Abbeel)

38 / 51

Problems with Value Iteration

- The 'max' value at each state rarely changes
- The policy often converges long before the values converge

Policy iteration is an alternative approach, which is still optimal and can converge much more quickly.

39 / 51

Policy iteration

$$\pi^{(0)} \xrightarrow{E} V^{\pi^{(0)}} \xrightarrow{I} \pi^{(1)} \xrightarrow{E} V^{\pi^{(1)}} \xrightarrow{I} \pi^{(2)} \xrightarrow{E} \dots$$

Repeat (until policy converges):

- Evaluate (E) V^{π} (where π is current policy)
- Policy improvement (I):

$$\pi'(s) \leftarrow \arg \max_{a \in A} \left[r(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V^{\pi}(s') \right], \quad \forall s$$

update policy using one-step look-ahead with V^{π} as future values

- $\pi \leftarrow \pi'$

Proof of convergence shows $V^{\pi^{(k+1)}} > V^{\pi^{(k)}}$ (if policy changes).

40 / 51

Example: Policy iteration

Example on a different grid world, initialized with $\pi(s) = \uparrow$ (all states).

0	0	0	1
0		0	-100
0	0	0	0

Z. Kolter

Original reward function

41 / 51

Example: Policy iteration

Example on a different grid world, initialized with $\pi(s) = \uparrow$ (all states).

0.418	0.884	2.331	6.367
0.367		-8.610	-105.7
-0.168	-4.641	-14.27	-85.05

Z. Kolter

V^π at one iteration

42 / 51

Example: Policy iteration

Example on a different grid world, initialized with $\pi(s) = \uparrow$ (all states).

5.414	6.248	7.116	8.634
4.753		2.881	-102.7
2.251	1.977	1.849	-8.701

Z. Kolter

V^π at two iterations

43 / 51

Example: Policy iteration

Example on a different grid world, initialized with $\pi(s) = \uparrow$ (all states).

5.470	6.313	7.190	8.669
4.803		3.347	-96.67
4.161	3.654	3.222	1.526

Z. Kolter

V^π at three iterations (converged!)

44 / 51

Typical Gridworld results

- Approximation of value function
 - Policy iteration: exact value function after three iterations
 - Value iteration: $\|\mathbf{V} - \mathbf{V}^*\|_2 < 10^{-4}$ after 100 iterations
- Approximation of optimal policy
 - Policy iteration: optimal policy after three iterations
 - Value iteration: optimal policy after 10 iterations

45 / 51

What is the difference?

Value iteration

$$V'(s) \leftarrow \max_{a \in A} \left[r(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V(s') \right], \quad \forall s$$

Policy iteration

$$\pi'(s) \leftarrow \arg \max_{a \in A} \left[r(s, a) + \gamma \sum_{s' \in S} p(s' | s, a) V^\pi(s') \right], \quad \forall s$$

46 / 51

Policy iteration or Value iteration?

Both converge to the optimal policy in a finite number of steps.

- Value iteration:

- $O(|S| \cdot |A| \cdot L)$ per iteration
- less work per iteration (no policy evaluation!)

- Policy iteration:

- policy changes every iteration
- $O(|S| \cdot |A| \cdot L + |S|^3)$ computation per iteration
- tends to require less steps (larger changes each step)

In practice, PI tends to be faster, especially if transition matrix is sparse so that policy evaluation is fast.

47 / 51

Other solution approaches

- Can take derivatives of a policy that is parameterized (good for large/continuous action spaces)
- Tree search: can “roll out,” or simulate policies. Good for large state spaces. (Approximate form of expectimax).
- Linear programming.

48 / 51

Contents

1 Introduction

2 Planning (finite horizon)


3 Planning (infinite horizon)

- Bellman equations
- Value Iteration
- Policy Iteration

4 Conclusion

49 / 51

Next Class: Learning a Policy



- Agent knows current state s takes actions a , and gets reward r .
- Only access to reward model $r(s, a)$, transition model $p(s' | s, a)$ via feedback
- Very challenging problem to learn π while uncertain about model of the world.

50 / 51

- MDPs are a general, probabilistic model for acting in an uncertain environment
- The main assumptions in the model are:
 - Markovian: $p_t(s_{t+1} \mid s_1, \dots, s_t, a_1, \dots, a_t) = p_t(s_{t+1} \mid s_t, a_t)$
 - Stationarity: $p_t(s_{t+1} \mid s_t, a_t) = p(s_{t+1} \mid s_t, a_t)$
- Planning is the problem of deciding how to act, given knowledge of the MDP (S, A, r, p)
- For the infinite time horizon, discounted setting, we can use value iteration and policy iteration.